

---

# Adversarial Training of Neural Networks against Systematic Uncertainty

---

Gilles Louppe  
New York University  
g.louppe@nyu.edu

## Abstract

## 1 Introduction

[GL: Distinction between statistical and systematic uncertainty.] [GL: Define nuisance parameters.]  
[GL: We want to build an accurate classifier whose output remains invariant with respect to systematic uncertainties.] [GL: Motivate the criterion (which may not be obvious for the ML crowd). See pivotal quantity motivation.]

## 2 Problem statement

Let assume a probability space  $(\Omega, \mathcal{F}, P)$ , where  $\Omega$  is a sample space,  $\mathcal{F}$  is a set of events and  $P$  is a probability measure. Let consider the multivariate random variables  $X_\lambda : \Omega \mapsto \mathbb{R}^p$  and  $Y : \Omega \mapsto \mathcal{Y}$ , where  $X_\lambda$  depends on a nuisance parameter  $\lambda \in \Lambda \subseteq \mathbb{R}$  whose values define a parameterized family of its systematic uncertainties. That is,  $X_\lambda$  and  $Y$  induce together a joint probability distribution  $p(X, Y|\lambda)$ , where the conditional on  $\lambda$  denotes  $X_\lambda$ . For training, let further assume a finite set  $\{x_i, y_i, \lambda_i\}_{i=1}^N$  of realizations  $X_{\lambda_i}(\omega_i), Y(\omega_i)$ , for  $\omega_i \in \Omega$  and known values  $\lambda_i$  of the nuisance parameter. Our goal is to learn a function  $f(\cdot; \theta_f) : \mathbb{R}^p \mapsto \mathcal{Y}$  of parameters  $\theta_f$  (e.g., a neural network-based classifier if  $\mathcal{Y}$  is a finite set of classes) and minimizing a loss  $\mathcal{L}_f(\theta_f)$ . In addition, we require that  $f(X_\lambda; \theta_f)$  should be robust to the value of the nuisance parameter  $\lambda$  – which remains unknown at test time. More specifically, we aim at building  $f$  such that in the ideal case

$$f(X_{\lambda_i}(\omega); \theta_f) = f(X_{\lambda_j}(\omega); \theta_f) \quad (1)$$

for any sample  $\omega \in \Omega$  and any  $\lambda_i, \lambda_j$  pair of values of the nuisance parameter.

Since we do not have training tuples  $(X_{\lambda_i}(\omega), X_{\lambda_j}(\omega))$  (for the same unknown  $\omega$ ), we propose instead to solve the closely related problem of finding a predictive function  $f$  such that

$$P(\{\omega | f(X_{\lambda_i}(\omega); \theta_f) = y\}) = P(\{\omega' | f(X_{\lambda_j}(\omega'); \theta_f) = y\}) \text{ for all } y \in \mathcal{Y}. \quad (2)$$

In words, we are looking for a predictive function  $f$  which is a pivotal quantity with respect to the nuisance parameter. That is, such that the distribution of  $f(X_\lambda; \theta_f)$  is invariant with respect to  $\lambda$ . Note that a function  $f$  for which Eqn. 1 is true necessarily satisfies Eqn. 2. The converse is however in general not true, since the sets of samples  $\{\omega | f(X_{\lambda_i}(\omega); \theta_f) = y\}$  and  $\{\omega' | f(X_{\lambda_j}(\omega'); \theta_f) = y\}$  do not need to be the same for the equality to hold.

## 3 Method

Adversarial training was first proposed by [1] as a way to build a generative model capable of producing samples from random noise  $z \sim p_Z$ . More specifically, the authors pit a generative model

$g : \mathcal{Z} \mapsto \mathbb{R}^p$  against an adversary classifier  $d : \mathbb{R}^p \mapsto \{0, 1\}$  whose antagonistic objective is to recognize real data  $X$  from generated data  $g(Z)$ . Both models  $g$  and  $d$  are trained simultaneously, in such a way that  $g$  learns to produce samples that are difficult to identify by  $d$ , while  $d$  incrementally adapts to changes in  $g$ . At the equilibrium,  $g$  models a distribution whose samples can be identified by  $d$  only by chance. That is, assuming enough capacity in  $d$  and  $g$ , the distribution  $p_{g(Z)}$  eventually converges towards the real distribution  $p_X$ .

In this work, we repurpose adversarial training as a means to regularize the predictive model  $f$  in order to satisfy Eqn. 2. In particular, we pit  $f$  against an adversary model  $r(\cdot; \theta_r) : \mathbb{R} \mapsto \Lambda$  of parameters  $\theta_r$  and associated loss  $\mathcal{L}_r(\theta_r)$ . This model takes as input realizations of  $f(X_\lambda; \theta_f)$  and produces as output predictions  $\hat{\lambda} = r(f(X_\lambda; \theta_f))$  of the nuisance parameter. If  $p(f(X_\lambda; \theta_f))$  varies with  $\lambda$ , then the corresponding correlation can be captured by  $r$ . By contrast, if  $p(f(X_\lambda; \theta_f))$  is invariant with  $\lambda$  as we require, then  $r$  should perform poorly. Training  $f$  such that it additionally minimizes the performance of  $r$  therefore acts as a regularization towards Eqn. 2.

As for generative adversarial networks, we propose to train  $f$  and  $r$  simultaneously, which we carry out by considering the value function

$$E(\theta_f, \theta_r) = \mathcal{L}_f(\theta_f) - \mathcal{L}_r(\theta_r) \quad (3)$$

that we optimize by finding the saddle point  $(\hat{\theta}_f, \hat{\theta}_r)$  such that

$$\hat{\theta}_f = \arg \min_{\theta_f} E(\theta_f, \hat{\theta}_r), \quad (4)$$

$$\hat{\theta}_r = \arg \max_{\theta_r} E(\hat{\theta}_f, \theta_r). \quad (5)$$

[GL: Expand the equation to make it clearer that  $f$  is plugged into  $r$ .] [GL: Explain this can be done simply using SGD.] [GL: What are the necessary conditions on  $\mathcal{L}_r$  to imply Eqn. 2 at the saddle point? We should clarify and prove that formally! There may also be assumptions required on the learning problem itself.]

## 4 Experiments

## 5 Related work

[GL: Similar to domain adaptation, but with infinitely many domains, as parameterized by  $\lambda$ , also related to transfer learning.]

## 6 Conclusions

### Acknowledgments

### References

- [1] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, “Generative adversarial nets,” in *Advances in Neural Information Processing Systems*, pp. 2672–2680, 2014.