Adversarial Training of Neural Networks against Systematic Uncertainty

Gilles Louppe
New York University
`g.louppe@nyu.edu`

Abstract

1   Introduction

[GL: Distinction between statistical and systematic uncertainty.] [GL: Define nuisance parameters.] [GL: We want to build an accurate classifier whose output remains invariant with respect to systematic uncertainties.] [GL: Motivate the criterion (which may not be obvious for the ML crowd). See pivotal quantity motivation.]

2   Problem statement

Let assume a probability space $(\Omega, \mathcal{F}, P)$, where Ω is a sample space, $\mathcal{F}$ is a set of events and P is a probability measure. Let consider the multivariate random variables $X_z : \Omega \mapsto \mathbb{R}^p$ and $Y : \Omega \mapsto \mathcal{Y}$, where X_z denotes a dependence on a nuisance parameter Z whose values $z \in \mathcal{Z}$ define a parameterized family of its systematic uncertainties. That is, X_z and Y induce together a joint probability distribution $p(X, Y | z)$, where the conditional on z denotes X_z. For training, let further assume a finite set $\{x_i, y_i, z_i\}_{i=1}^N$ of realizations $X_{z_i}(\omega_i), Y(\omega_i)$, for $\omega_i \in \Omega$ and known values z_i of the nuisance parameter. Our goal is to learn a function $f(\cdot; \theta_f) : \mathbb{R}^p \mapsto \mathcal{Y}$ of parameters θ_f (e.g., a neural network-based classifier if $\mathcal{Y}$ is a finite set of classes) and minimizing a loss $\mathcal{L}_f(\theta_f)$. In addition, we require that $f(X_z; \theta_f)$ should be robust to the value z of the nuisance parameter – which remains unknown at test time. More specifically, we aim at building f such that in the ideal case

$$f(X_z(\omega); \theta_f) = f(X_{z'}(\omega); \theta_f) \tag{1}$$

for all samples $\omega \in \Omega$ and all z, z' pairs of values of the nuisance parameter.

Since we do not have training tuples $(X_z(\omega), X_{z'}(\omega))$ (for the same unknown ω), we propose instead to solve the closely related problem of finding a predictive function f such that

$$P(\{\omega | f(X_z(\omega); \theta_f) = y\}) = P(\{\omega' | f(X_{z'}(\omega'); \theta_f) = y\}) \text{ for all } y \in \mathcal{Y}. \tag{2}$$

In words, we are looking for a predictive function f which is a pivotal quantity [1] with respect to the nuisance parameter. That is, such that the distribution of $f(X_z; \theta_f)$ is invariant with respect to the value z of the nuisance. Note that a function f for which Eqn. 1 is true necessarily satisfies Eqn. 2. The converse is however in general not true, since the sets of samples $\{\omega | f(X_z(\omega); \theta_f) = y\}$ and $\{\omega' | f(X_{z'}(\omega'); \theta_f) = y\}$ do not need to be the same for the equality to hold. In order to simplify notations, and as only Eqn. 2 is of direct interest in this work, we denote from here on the pivotal quantity criterion as

$$p(f(X; \theta_f) | z) = p(f(X; \theta_f) | z') \text{ for all } z, z' \in \mathcal{Z}. \tag{3}$$

3 Method

Adversarial training was first proposed by [2] as a way to build a generative model capable of producing samples from random noise $z \sim p_Z$. More specifically, the authors pit a generative model $g : \mathcal{Z} \mapsto \mathbb{R}^p$ against an adversary classifier $d : \mathbb{R}^p \mapsto \{0, 1\}$ whose antagonistic objective is to recognize real data X from generated data $g(Z)$. Both models g and d are trained simultaneously, in such a way that g learns to produce samples that are difficult to identify by d, while d incrementally adapts to changes in g. At the equilibrium, g models a distribution whose samples can be identified by d only by chance. That is, assuming enough capacity in d and g, the distribution $p_{g(Z)}$ eventually converges towards the real distribution p_X.

In this work, we repurpose adversarial training as a means to constraint the predictive model f in order to satisfy Eqn. 3. In particular, we pit f against an adversary classifier $r(\cdot; \theta_r) : \mathbb{R} \mapsto \mathcal{Z}$ of parameters θ_r and associated loss $\mathcal{L}_r(\theta_f, \theta_r)$. Assuming that $\mathcal{Z}$ defines a finite family of nuisance values z_l (for $l = 1, \ldots, |\mathcal{Z}|$), this classifier takes as input realizations of $f(X; \theta_f)$, for the current value θ_f of f parameters, and produces as output probability estimates $r(f(X; \theta_f); \theta_r)_l = \hat{p}(z_l | f(X; \theta_f))$ that $f(X; \theta_f)$ is generated from the nuisance value z_l. If $p(f(X; \theta_f)|z)$ varies with z, then the corresponding correlation can be captured by r. By contrast, if $p(f(X; \theta_f)|z)$ is invariant with z, as we require, then r should perform poorly and be close to random guessing. Training f such that it additionally minimizes the performance of r therefore acts as a regularization towards Eqn. 3.

As for generative adversarial networks, we propose to train f and r simultaneously, which we carry out by considering the value function

$$E(\theta_f, \theta_r) = \mathcal{L}_f(\theta_f) - \mathcal{L}_r(\theta_f, \theta_r) \tag{4}$$

that we optimize by finding the saddle point $(\hat{\theta}_f, \hat{\theta}_r)$ such that

$$\hat{\theta}_f = \arg\min_{\theta_f} E(\theta_f, \hat{\theta}_r), \tag{5}$$

$$\hat{\theta}_r = \arg\max_{\theta_r} E(\hat{\theta}_f, \theta_r). \tag{6}$$

The adversarial training procedure to obtain $(\hat{\theta}_f, \hat{\theta}_r)$ is formally presented in Algorithm 1 in the case of f being a classifier and of the cross-entropy loss for both $\mathcal{L}_f$ and $\mathcal{L}_r$. The algorithm consists in using stochastic gradient descent alternatively to optimize Eqn. 5 and 6.

Algorithm 1 Adversarial training of a classifier f against an adversary r.

Inputs: training data $\{x_i, y_i, z_i\}_{i=1}^N$
Outputs: $\hat{\theta}_f, \hat{\theta}_r$
Hyper-parameters: Number T of training iterations, Number K of gradient steps to update r.

1: **for** $t = 1$ to T **do**
2:     **for** $k = 1$ to K **do** $\triangleright$ Update r
3:         Sample minibatch $\{x_m, z_m\}_{m=1}^M$ of size M;
4:         With θ_f fixed, update r by ascending its stochastic gradient $\nabla_{\theta_r} E(\theta_f, \theta_r) :=$

$$\nabla_{\theta_r} \sum_{m=1}^M \left[\sum_{z_l \in \mathcal{Z}} 1(z_m = z_l) \log r(f(x_m; \theta_f); \theta_r)_l \right];$$

5:     **end for**
6:     Sample minibatch $\{x_m, y_m, z_m\}_{m=1}^M$ of size M; $\triangleright$ Update f
7:     With θ_r fixed, update f by descending its stochastic gradient $\nabla_{\theta_f} E(\theta_f, \theta_r) :=$

$$\nabla_{\theta_f} \sum_{m=1}^M \left[- \sum_{y_c \in \mathcal{Y}} 1(y_m = y_c) \log f(x_m; \theta_f)_c + \sum_{z_l \in \mathcal{Z}} 1(z_m = z_l) \log r(f(x_m; \theta_f); \theta_r)_l \right];$$

8: **end for**

---

4 Theoretical results

In this section, we show that in the setting of Algorithm 1, the procedure converges to a classifier f which is a pivotal quantity in the sense of Eqn. 3. Results below are derived in a non-parametric setting, by assuming that both f and r have enough capacity. To simplify the presentation, we also assume the uniform prior $p(z) = \frac{1}{|\mathcal{Z}|}$ for all $z \in \mathcal{Z}$, e.g. by having the same number of training samples for each modality z of the nuisance parameter.

Proposition 1. *Let θ_f be fixed and $\hat{\theta}_r = \arg\max_{\theta_r} E(\theta_f, \theta_r)$. If $r(f(X; \theta_f); \hat{\theta}_r)_l = \frac{1}{|\mathcal{Z}|}$ for all z_l, then f is a pivotal quantity.*

Proof. Let us first recall that the cross-entropy for distributions p and q is minimized when $p = q$. For $\mathcal{L}_r$ defined as the cross-entropy between the true conditional distribution of the nuisance $p_{Z|f(X;\theta_f)}$ and the approximate conditional distribution of the nuisance $p_{r(f(X;\theta_f);\theta_r)|f(X)}$, the optimal parameters $\hat{\theta}_r = \arg\max_{\theta_r} E(\theta_f, \theta_r) = \arg\min_{\theta_r} L_r(\theta_f, \theta_r)$ are therefore such that $p_{r(f(X;\theta_f);\hat{\theta}_r)|f(X)} = p_{Z|f(X;\theta_f)}$.

In other words, for all $z_l \in \mathcal{Z}$, we have $r(f(X; \theta_f); \hat{\theta}_r)_l = p(z_l|f(X;\theta_f))$. By assumption, $r(f(X;\theta_f); \hat{\theta}_r)_l = \frac{1}{|\mathcal{Z}|}$, and therefore $p(z_l|f(X;\theta_f)) = \frac{1}{|\mathcal{Z}|}$. Using the Bayes' rule, we write

$$
\begin{aligned}
p(f(X;\theta_f)|z_l) &= \frac{p(z_l|f(X;\theta_f))p(f(X;\theta_f))}{p(z_l)} \\
&= \frac{\frac{1}{|\mathcal{Z}|}p(f(X;\theta_f))}{\frac{1}{|\mathcal{Z}|}} \\
&= p(f(X;\theta_f)),
\end{aligned}
$$

which holds for all $z_l \in \mathcal{Z}$ and implies that f is a pivotal quantity. $\qquad\square$

Proposition 2. *If there exists a saddle point $(\hat{\theta}_f, \hat{\theta}_r)$ for Eqn. 5 and 6 such that $E(\hat{\theta}_f, \hat{\theta}_r) = H(p_{Y|X}) - \log|\mathcal{Z}|$, then $f(\cdot; \hat{\theta}_f)$ is both an optimal classifier with respect to L_f and a pivotal quantity.*

Proof. For fixed θ_f, the adversary r is optimal at $\hat{\theta}_r = \arg\max_{\theta_r} E(\theta_f, \theta_r) = \arg\min_{\theta_r} L_r(\theta_f, \theta_r)$, in which case $p_{r(f(X;\theta_f);\hat{\theta}_r)|f(X)} = p_{Z|f(X;\theta_f)}$ and L_r reduces to the entropy $H(p_{Z|f(X;\theta_f)})$ of the conditional distribution of the nuisance. The value function E can therefore be rewritten as

$$
E'(\theta_f) = L_f(\theta_f) - H(p_{Z|f(X;\theta_f)}).
$$

In particular, we have the lower bound $H(p_{Y|X}) - \log|\mathcal{Z}| \leq L_f(\theta_f) - H(p_{Z|f(X;\theta_f)})$ where the equality holds at $\hat{\theta}_f = \arg\min_{\theta_f} E'(\theta_f)$ only when

- $\hat{\theta}_f$ corresponds to the parameters of an optimal classifier, in which case the log-loss L_f reduces to $H(p_{Y|X})$,

- all outcomes of $Z|f(X; \hat{\theta}_f)$ are equally likely, in which case $p(z_l|f(X; \hat{\theta}_f)) = \frac{1}{|\mathcal{Z}|}$ for all $z_l \in \mathcal{Z}$ and $H(p_{Z|f(X;\hat{\theta}_f)}) = -\sum_{z_l \in \mathcal{Z}} p(z|\hat{\theta}_f) \log p(z|\hat{\theta}_f) = -\sum_{z_l \in \mathcal{Z}} \frac{1}{|\mathcal{Z}|} \log \frac{1}{|\mathcal{Z}|} = \log|\mathcal{Z}|$.

Accordingly, the second condition implies that $r(f(X; \hat{\theta}_f); \hat{\theta}_r)_l = \frac{1}{|\mathcal{Z}|}$ and therefore that at this point, because of Proposition 1, the optimal classifier $f(\cdot; \hat{\theta}_f)$ is also a pivotal quantity. $\qquad\square$

[GL: We should further discuss that in practice, the equality may never hold. We should discuss in which circumstances. In such case, the pivotal quantity constraint can however be enforced by outweighting the L_r, resulting in a trade-off between classifier optimality and pivotality.]

Proposition 3. *[GL: It remains to prove that the procedure of Algorithm 1 converges towards that saddle point. The proof should be similar to the proof of convergence in the GAN paper.]*

5 Experiments

6 Related work

[GL: Similar to domain adaptation, but with infinitely many domains, as parameterized by Z, also related to transfer learning.]

7 Conclusions

Acknowledgments

References

[1] M. H. Degroot and M. J. Schervish, *Probability and statistics*. 4 ed., 2010.

[2] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative adversarial nets," in *Advances in Neural Information Processing Systems*, pp. 2672–2680, 2014.